

## Contents

1.Introduction .....	2
2.Initial Network Configuration.....	2
2.1.Proof of Concept .....	2
2.2.OpenFlow Packet Analysis .....	5
3.Network Functionalities.....	10
3.1.Firewall Implementation .....	10
3.2.Load Balancing.....	11
3.3.Network Monitoring .....	16

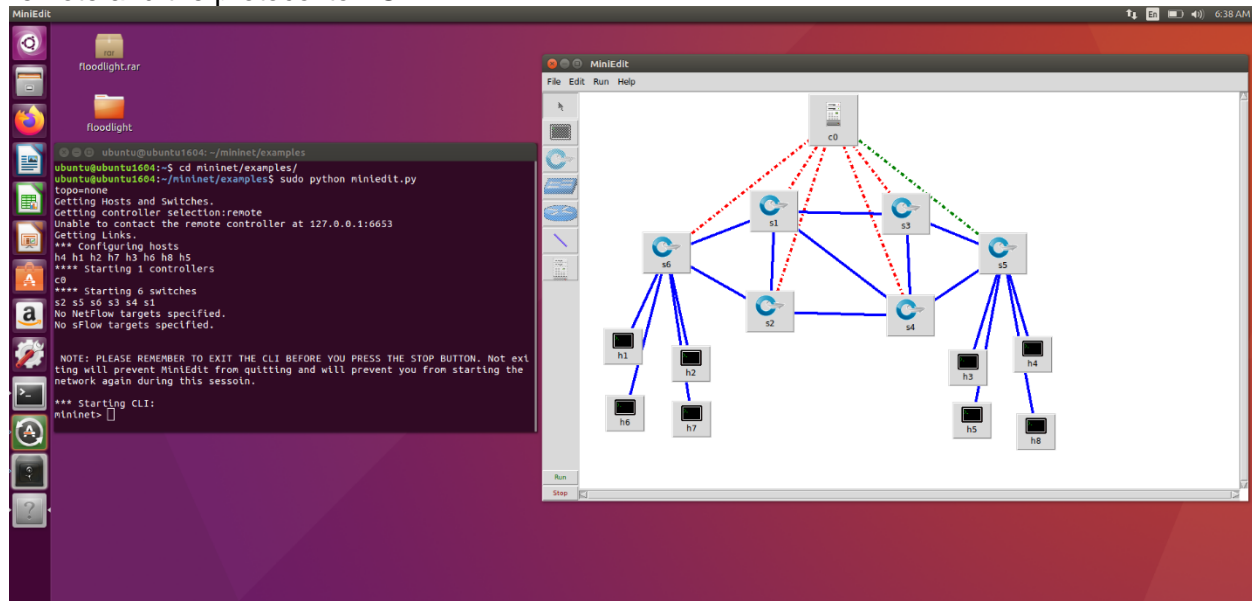
## 1. Introduction

The purpose of this report is to analyse and examine a network infrastructure using SDN core concepts and technologies. This will include a proof-of-concept design using an open source SDN solution known as Mininet, as well as an overview of the Floodlight Openflow controller, network ping tests, packet analysis with Wireshark, firewall implementation, load balancing and network monitoring.

## 2. Initial Network Configuration

### 2.1. Proof of Concept

The topology was created via Miniedit, accessed through a command line python script as shown in figure 1. The “c0” represents the central network floodlight controller, which is connected to 6 switches, delegating traffic between 8 different hosts machines. Additional options that were required to be changed were the controller port to 6653, the controller type to remote and the protocol to TCP.



After the topology was created, the Floodlight Openflow SDN controller was implemented and chosen. Figure 2 shows the hosts connecting to the switches IPs using the TCP 3-way handshake process

Figure 1: Network Topology

```
ubuntu@ubuntu1604: ~/Desktop/floodlight
2022-03-21 07:14:14.334 INFO [n.f.f.Forwarding] Default flags will be set to SEND_FLOW_REM false.
2022-03-21 07:14:14.334 INFO [n.f.f.Forwarding] default flow matches set to: IN_PORT=true, VLAN=true, MAC=true, IP=true, FLAG=true, TPPT=true
2022-03-21 07:14:14.335 INFO [n.f.f.Forwarding] Default detailed flow matches set to: SRC_MAC=true, DST_MAC=true, SRC_IP=true, DST_IP=true, SRC_TPPT=true, DST_TPPT=true
2022-03-21 07:14:14.335 INFO [n.f.f.Forwarding] Not flooding ARP packets. ARP flows will be inserted for known destinations
2022-03-21 07:14:14.336 INFO [n.f.f.Forwarding] Flows will be removed on link/port down events
2022-03-21 07:14:14.336 INFO [n.f.s.StatisticsCollector] Statistics collection disabled
2022-03-21 07:14:14.336 INFO [n.f.s.StatisticsCollector] Port statistics collection interval set to 10s
2022-03-21 07:14:14.598 INFO [o.s.s.l.SyncManager] [1] Updating sync configuration ClusterConfig [allNodes=[1-Node [hostname=192.168.1.100, port=6642, nodeId=1, domainId=1], 2-Node [hostname=192.168.1.100, port=6643, nodeId=2, domainId=1]], autoScheme=CHALLENGE_RESPONSE, keyStorePath=/etc/floodlight/keys/keys, keyStorePassword is set]
2022-03-21 07:14:15.919 INFO [o.s.s.l.RPCService] Listening for internal floodlight RPC on 0.0.0.0/0.0.0.0:6642
2022-03-21 07:14:15.909 INFO [o.r.c.i.Server] Starting the Simple [HTTP/1.1] server on port 8080
2022-03-21 07:14:15.910 INFO [o.r.restlet] Starting net.floodlightcontroller.restserver.RestAppServer$RestApplication application
2022-03-21 07:14:16.238 INFO [n.f.c.l.OFChannelHandler] New switch connection from /127.0.0.1:45280
2022-03-21 07:14:16.239 INFO [n.f.c.l.OFChannelHandler] New switch connection from /127.0.0.1:45286
2022-03-21 07:14:16.239 INFO [n.f.c.l.OFChannelHandler] New switch connection from /127.0.0.1:45276
2022-03-21 07:14:16.238 INFO [n.f.c.l.OFChannelHandler] New switch connection from /127.0.0.1:45282
2022-03-21 07:14:16.238 INFO [n.f.c.l.OFChannelHandler] New switch connection from /127.0.0.1:45284
2022-03-21 07:14:16.238 INFO [n.f.c.l.OFChannelHandler] New switch connection from /127.0.0.1:45278
2022-03-21 07:14:16.445 INFO [n.f.c.l.OFChannelHandler] Negotiated OpenFlow version of OF_13 for /127.0.0.1:45280 with greatest version bitmap algorithm.
2022-03-21 07:14:16.446 INFO [n.f.c.l.OFChannelHandler] Negotiated OpenFlow version of OF_13 for /127.0.0.1:45278 with greatest version bitmap algorithm.
2022-03-21 07:14:16.446 INFO [n.f.c.l.OFChannelHandler] Negotiated OpenFlow version of OF_13 for /127.0.0.1:45276 with greatest version bitmap algorithm.
2022-03-21 07:14:16.446 INFO [n.f.c.l.OFChannelHandler] Negotiated OpenFlow version of OF_13 for /127.0.0.1:45282 with greatest version bitmap algorithm.
2022-03-21 07:14:16.445 INFO [n.f.c.l.OFChannelHandler] Negotiated OpenFlow version of OF_13 for /127.0.0.1:45284 with greatest version bitmap algorithm.
2022-03-21 07:14:16.459 INFO [n.f.c.l.OFChannelHandler] Negotiated OpenFlow version of OF_13 for /127.0.0.1:45278 with greatest version bitmap algorithm.
2022-03-21 07:14:16.598 INFO [n.f.c.l.OFChannelHandler] Switch OFSwitch DPID[00:00:00:00:00:00:03] bound to class class net.floodlightcontroller.core.internal.OFSwitch, description SwitchDescr
ption [manufacturerDescription=Nicira, Inc., hardwareDescription=Open vSwitch, softwareDescription=2.5.5, serialNumber=None, datapathDescription=None]
2022-03-21 07:14:16.599 INFO [n.f.c.l.OFChannelHandler] Switch OFSwitch DPID[00:00:00:00:00:00:06] bound to class class net.floodlightcontroller.core.internal.OFSwitch, description SwitchDescr
ption [manufacturerDescription=Nicira, Inc., hardwareDescription=Open vSwitch, softwareDescription=2.5.5, serialNumber=None, datapathDescription=None]
2022-03-21 07:14:16.599 INFO [n.f.c.l.OFChannelHandler] Switch OFSwitch DPID[00:00:00:00:00:00:04] bound to class class net.floodlightcontroller.core.internal.OFSwitch, description SwitchDescr
ption [manufacturerDescription=Nicira, Inc., hardwareDescription=Open vSwitch, softwareDescription=2.5.5, serialNumber=None, datapathDescription=None]
2022-03-21 07:14:16.599 INFO [n.f.c.l.OFChannelHandler] Switch OFSwitch DPID[00:00:00:00:00:00:01] bound to class class net.floodlightcontroller.core.internal.OFSwitch, description SwitchDescr
ption [manufacturerDescription=Nicira, Inc., hardwareDescription=Open vSwitch, softwareDescription=2.5.5, serialNumber=None, datapathDescription=None]
2022-03-21 07:14:16.599 INFO [n.f.c.l.OFChannelHandler] Switch OFSwitch DPID[00:00:00:00:00:00:05] bound to class class net.floodlightcontroller.core.internal.OFSwitch, description SwitchDescr
ption [manufacturerDescription=Nicira, Inc., hardwareDescription=Open vSwitch, softwareDescription=2.5.5, serialNumber=None, datapathDescription=None]
2022-03-21 07:14:16.599 INFO [n.f.c.l.OFChannelHandler] Switch OFSwitch DPID[00:00:00:00:00:00:02] bound to class class net.floodlightcontroller.core.internal.OFSwitch, description SwitchDescr
ption [manufacturerDescription=Nicira, Inc., hardwareDescription=Open vSwitch, softwareDescription=2.5.5, serialNumber=None, datapathDescription=None]
2022-03-21 07:14:19.477 INFO [n.f.c.l.OFChannelHandler] Defining switch role from config file: ROLE_MASTER
2022-03-21 07:14:19.478 INFO [n.f.c.l.OFChannelHandler] Defining switch role from config file: ROLE_MASTER
2022-03-21 07:14:19.478 INFO [n.f.c.l.OFChannelHandler] Defining switch role from config file: ROLE_MASTER
2022-03-21 07:14:19.480 INFO [n.f.c.l.OFChannelHandler] Defining switch role from config file: ROLE_MASTER
2022-03-21 07:14:19.520 INFO [n.f.c.l.OFChannelHandler] Clearing flow tables of 00:00:00:00:00:00:03 on upcoming transition to MASTER.
2022-03-21 07:14:19.521 INFO [n.f.c.l.OFChannelHandler] Clearing flow tables of 00:00:00:00:00:00:01 on upcoming transition to MASTER.
2022-03-21 07:14:19.521 INFO [n.f.c.l.OFChannelHandler] Clearing flow tables of 00:00:00:00:00:00:04 on upcoming transition to MASTER.
2022-03-21 07:14:19.521 INFO [n.f.c.l.OFChannelHandler] Clearing flow tables of 00:00:00:00:00:00:05 on upcoming transition to MASTER.
2022-03-21 07:14:19.521 INFO [n.f.c.l.OFChannelHandler] Clearing flow tables of 00:00:00:00:00:00:02 on upcoming transition to MASTER.
2022-03-21 07:14:19.559 INFO [n.f.c.l.OFChannelHandler] Clearing flow tables of 00:00:00:00:00:00:03 on upcoming transition to MASTER.
2022-03-21 07:14:19.559 INFO [n.f.c.l.OFChannelHandler] Clearing flow tables of 00:00:00:00:00:00:01 on upcoming transition to MASTER.
2022-03-21 07:14:19.571 INFO [n.f.c.l.OFChannelHandler] Clearing flow tables of 00:00:00:00:00:00:04 on upcoming transition to MASTER.
2022-03-21 07:14:19.571 INFO [n.f.c.l.OFChannelHandler] Clearing flow tables of 00:00:00:00:00:00:05 on upcoming transition to MASTER.
2022-03-21 07:14:19.898 ERROR [n.f.c.l.OFChannelHandler] OFGroupDeletedErrorsVer13(xid=76, code=INVALID_GROUP, data=OFGroupDeleteVer13(xid=76, groupType=INDIRECT, group=all, buckets=[])) from switch OFSwitch DPID[00:00:00:00:00:00:01] in state net.floodlightcontroller.core.internal.OFSwitchHandshakeHandler$MasterState@53e5e
```

Figure 2: Topology Connections

Figure 3: Floodlight GUI

In contrast, another popular controller is OpenDaylight, which is similar to Floodlight and can provide northbound services to multiple stacks and of neutron networks. However, Floodlight provides more efficient services in terms of packet loss and latency.

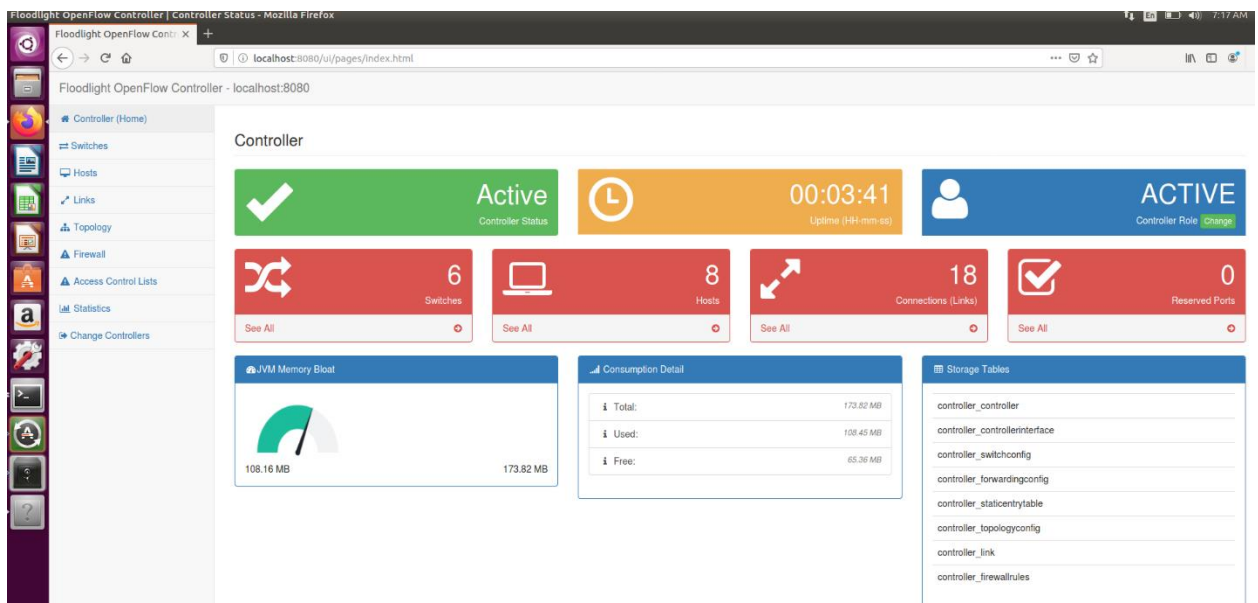






Figure 5: Topology Flow Table

Figure 3 shows the Floodlight GUI which has a visual representation of all the hosts and switches, and to test the network works, a ping all test was done in Mininet which resulted in 0 packet loss, showing the topology working correctly (Figure 4).

Figure 5 shows the topology flow table, meaning any packet information is displayed here and this further shows the network working correctly without any errors or drops.

```

root@ubuntu1604:~/mininet/examples# ping 10.0.0.2
PING 10.0.0.2 (10.0.0.2) 56(84) bytes of data:
64 bytes from 10.0.0.2: icmp_seq=1 ttl=64 time=0.89 ms
64 bytes from 10.0.0.2: icmp_seq=2 ttl=64 time=0.571 ms
64 bytes from 10.0.0.2: icmp_seq=3 ttl=64 time=0.084 ms
64 bytes from 10.0.0.2: icmp_seq=4 ttl=64 time=0.081 ms
64 bytes from 10.0.0.2: icmp_seq=5 ttl=64 time=0.094 ms
64 bytes from 10.0.0.2: icmp_seq=6 ttl=64 time=0.046 ms
64 bytes from 10.0.0.2: icmp_seq=7 ttl=64 time=0.044 ms
64 bytes from 10.0.0.2: icmp_seq=8 ttl=64 time=0.131 ms
64 bytes from 10.0.0.2: icmp_seq=9 ttl=64 time=0.077 ms
64 bytes from 10.0.0.2: icmp_seq=10 ttl=64 time=0.083 ms
64 bytes from 10.0.0.2: icmp_seq=11 ttl=64 time=0.056 ms
^C
--- 10.0.0.2 ping statistics ---
11 packets transmitted, 11 received, 0% packet loss, time 10200ms
rtt min/avg/max/mdev = 0.044/1.024/9.882/2.809 ms
root@ubuntu1604:~/mininet/examples#

ubuntu@ubuntu1604:~$ sudo ovs-ofctl dump-flows s6 -O OpenFlow13
OFPST_FLOW reply (OF1.3) (xid=0x2):
 cookie=0x200001770000000, duration=3.181s, table=0, n_packets=2, n_bytes=196, id
le_timeout=5, priority=1,ip,in_port=1,dl_src=fe:54:de:b6:17:4c,dl_dst=ce:5f:91:0
e:e2:51,nw_src=10.0.0.1,nw_dst=10.0.0.2 actions=output:2
 cookie=0x200001780000000, duration=3.180s, table=0, n_packets=2, n_bytes=196, id
le_timeout=5, priority=1,ip,in_port=2,dl_src=ce:5f:91:0e:e2:51,dl_dst=fe:54:de:b
6:17:4c,nw_src=10.0.0.2,nw_dst=10.0.0.1 actions=output:1
 cookie=0x0, duration=836.441s, table=0, n_packets=415, n_bytes=28934, priority=
0 actions=CONTROLLER:65535
ubuntu@ubuntu1604:~$
  
```

Figure 6: OpenFlow Table Test

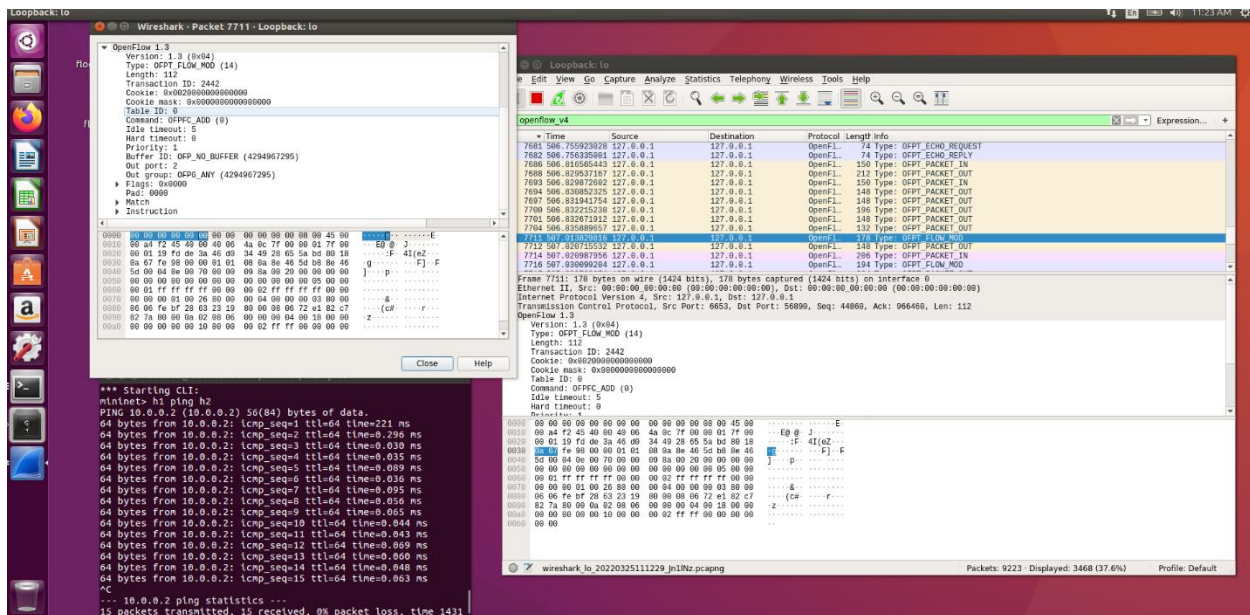
Figure 6 represents the test done to achieve the flow tables, and that each individual node can ping each other, concluding the network is setup accurately with no packet loss.

## 2.2. OpenFlow Packet Analysis

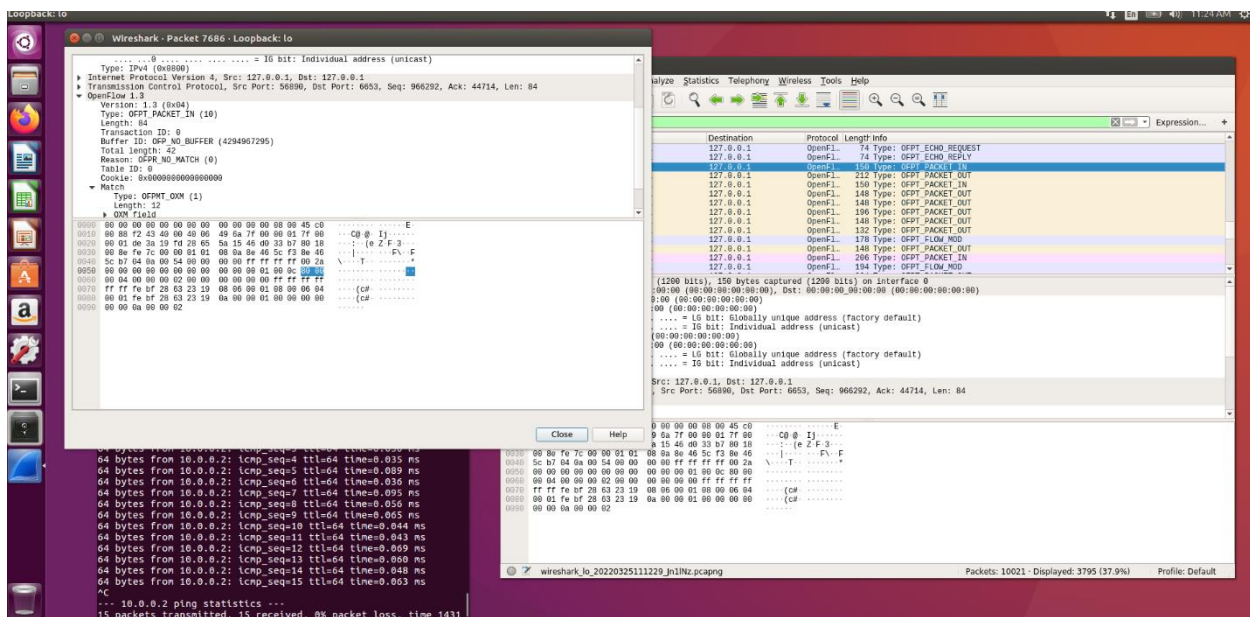
The Openflow packet analysis was achieved using Wireshark, where the Floodlight controller was started, and traffic was filtered to Loopback via openflow\_v4. This is because the network uses the Openflow framework and therefore this filter achieves relevant traffic (figure 7).

OFPT\_Hello establishes the handshake when an Openflow connection is confirmed and contains elements to help the connection setup. The Floodlight controller and switch uses OFPT\_Hello to also establish a familiar version number in order for messages to be communicated across the network.





Flow table modifications and configuration entries can be adjusted through the FLOW\_MOD traffic request. This is carried out from the controller, and therefore any new entries to the specified flow table are through this request (Figure 10).





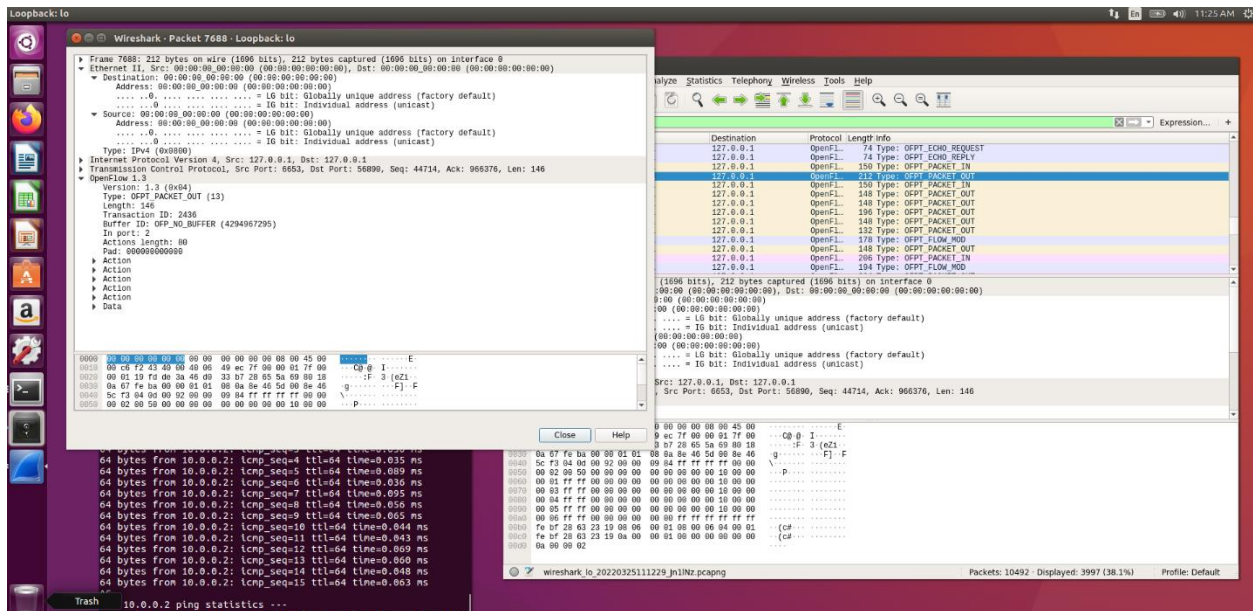


Figure 12: Wireshark OFPT\_Packet\_Out

Figures 11 and 12 show the Openflow packet in and out messages. The "packet in" messages use the data path to show packets received, which are then sent to the controller. The "packet out" message is used when the controller sends a packet back through the data path.

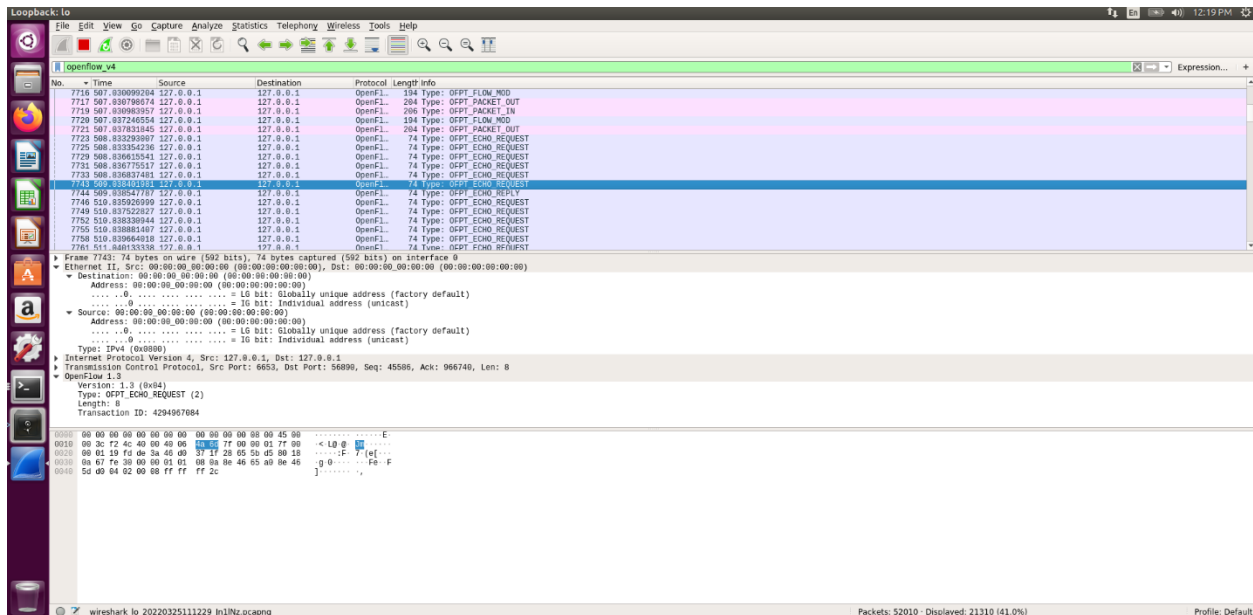


Figure 13: Wireshark OFPT\_Echo\_Request

Figure 13 shows the different types of echo-reply and sent messages across Wireshark. This means that the switch and controller can confirm an integrity between the connectivity of Openflow, to ensure connection is officially established. Each of these will help with connectivity troubleshooting and testing.

### 3. Network Functionalities

#### 3.1. Firewall Implementation

The firewall security policy that will be applied will be done through the access control list (ACL) as shown in Figure 14. Within the network topology (figure 1), the hosts are split between different switches. These hosts could represent different departments and it's possible to block specific traffic between these departments. Therefore, the ACL rule was created was to block ICMP traffic from the source, host 1, to the destination, host 8.

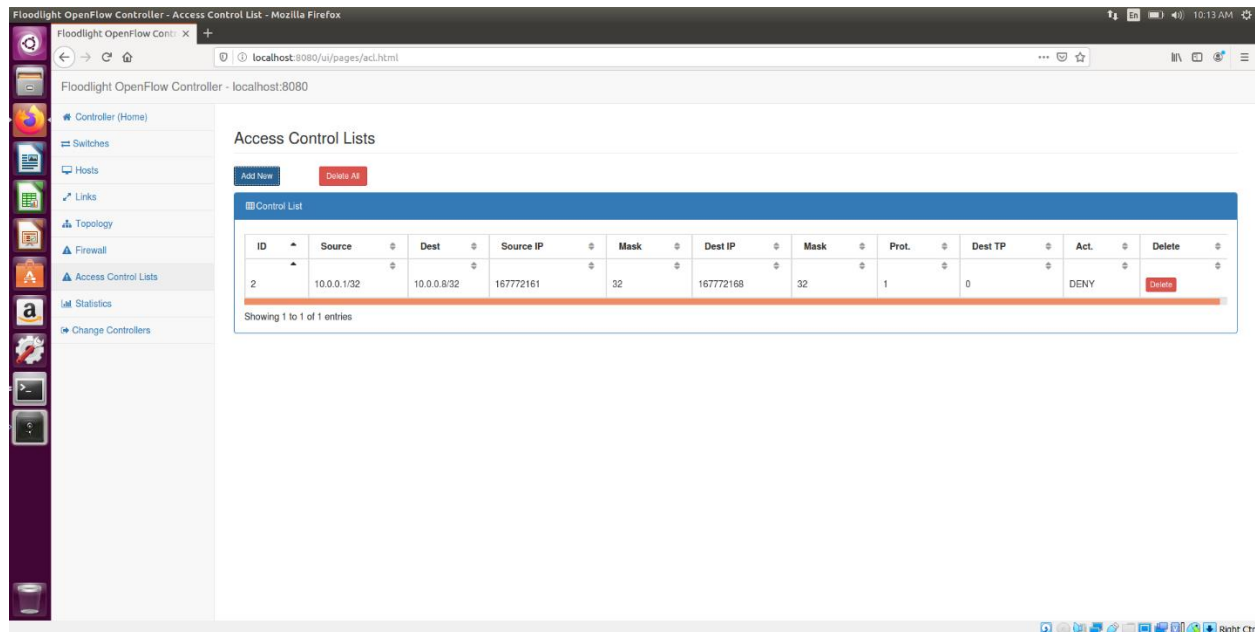


Figure 14: ACL

This is significant because without ACL firewall policies, it would be very easy to one department to access other departments data, and that the ports within the network would be open to the whole company, creating a security risk based on the principle of least privilege. Additionally, if one host was to become hijacked, that potential attacker could access network traffic within the whole network.

Alternatively, a company could implement a firewall to improve security posture of the network. This is because firewalls perform stateful inspection, resulting in each packet being inspected per flow, whereas ACL are limited to stateless which is per packet. Furthermore, ACL do not include intrusion detection and generally only include layer 4, however, ACL is very easy to setup to create a blocked connection between 2 hosts.

```
ubuntu@ubuntu1604: ~/mininet/examples
*** Starting controllers
*** Starting switches
*** Post configure switches and hosts
*** Starting CLI:
mininet> pingall
*** Ping: testing ping reachability
h7 -> h5 h2 h8 h1 h6 h4 h3
h5 -> h7 h2 h8 h1 h6 h4 h3
h2 -> h7 h5 h8 h1 h6 h4 X
h8 -> h7 h5 X X h6 h4 h3
h1 -> h7 h5 h2 h8 X h4 h3
h6 -> h7 h5 h2 h8 X h4 h3
h4 -> h7 h5 h2 h8 h1 h6 h3
h3 -> h7 h5 h2 h8 h1 h6 h4
*** Results: 8% dropped (51/56 received)
mininet> pingall
*** Ping: testing ping reachability
h7 -> h5 h2 h8 h1 h6 h4 h3
h5 -> h7 h2 h8 h1 h6 h4 h3
h2 -> h7 h5 h8 h1 h6 h4 h3
h8 -> h7 h5 h2 X h6 h4 h3
h1 -> h7 h5 h2 X h6 h4 h3
h6 -> h7 h5 h2 h8 h1 h4 h3
h4 -> h7 h5 h2 h8 h1 h6 h3
h3 -> h7 h5 h2 h8 h1 h6 h4
*** Results: 3% dropped (54/56 received)
mininet> 
```

Figure 15: PingAll Test

Consequently, as a result of the ACL rule implementation between hosts 1 and 8, a Pingall test was carried out, as shown in Figure 15. This is to ensure that the ACL filter rule works correctly within the network. An incorrect test would mean that the ACL rule is working incorrectly, and therefore could be a security issue.

Using the command line within Ubuntu, the Pingall test within the Mininet interface resulted in only 2 packets dropped (figure 15), and those packets dropped are represented by the 'X', shown on host 1 to host 8, and on host 8 to host one, resulting in a successful ACL implementation.

### 3.2. Load Balancing

In order to balance network loads across distributed paths, load balancing needs to be introduced to the network. A company with 250 employees simultaneously using a small network topology would need the most efficient distributed data paths, due to high traffic areas being managed and capacity utilisation being upheld to ensure one server is not overworked.

Additionally, load balancing performs different functions such as only sending traffic to online servers to ensure availability and reliability of requests, which can occur from both client and network side of things. For example, request happen either end of the network, from the controller, switches, or hosts. Not all 250 employees will be logged on simultaneously, therefore, if all the servers were in use, capacity utilisation will be wasted, incurring higher costs.

However, load balancing provides a flexible service where servers can be added or subtracted anytime depending on demand levels, meaning expenses and capital of running the servers will decrease. (Creating manually – installing flows in flow table of switch)

```

ubuntu@ubuntu1604: ~/mininet/examples
*** Results: 8% dropped (51/56 received)
mininet> pingall
*** Ping: testing ping reachability
h7 -> h5 h2 h8 h1 h6 h4 h3
h5 -> h7 h2 h8 h1 h6 h4 h3
h2 -> h7 h5 h8 h1 h6 h4 h3
h8 -> h7 h5 h2 X h6 h4 h3
h1 -> h7 h5 h2 X h6 h4 h3
h6 -> h7 h5 h2 h8 h1 h4 h3
h4 -> h7 h5 h2 h8 h1 h6 h3
h3 -> h7 h5 h2 h8 h1 h6 h4
*** Results: 3% dropped (54/56 received)
mininet> net
h7 h7-eth0:s6-eth1
h5 h5-eth0:s5-eth6
h2 h2-eth0:s6-eth3
h8 h8-eth0:s5-eth1
h1 h1-eth0:s6-eth2
h6 h6-eth0:s6-eth4
h4 h4-eth0:s5-eth5
h3 h3-eth0:s5-eth4
s1 lo: s1-eth1:s6-eth5 s1-eth2:s3-eth1
s4 lo: s4-eth1:s2-eth2 s4-eth2:s5-eth2
s3 lo: s3-eth1:s1-eth2 s3-eth2:s5-eth3
s6 lo: s6-eth1:h7-eth0 s6-eth2:h1-eth0 s6-eth3:h2-eth0 s6-eth4:h6-eth0 s6-eth5:s1-eth1 s6-eth6:s2-eth1
s5 lo: s5-eth1:h8-eth0 s5-eth2:s4-eth2 s5-eth3:s3-eth2 s5-eth4:h3-eth0 s5-eth5:h4-eth0 s5-eth6:h5-eth0
s2 lo: s2-eth1:s6-eth6 s2-eth2:s4-eth1
c0
mininet>

```

Figure 16: Mininet Network

Figure 17: Adding Flows

Figure 16 shows a connection peered Mininet network, which is an overview from the topology as shown in figure 1. From this, connections can be identified in order to make load balancing changes to the network. For

```

ubuntu@ubuntu1604:~$ sudo ovs-ofctl -O OpenFlow13 add-flow s6 in_port=1,priority
=40000,dl_type=0x800,nw_src=10.0.0.1,nw_dst=10.0.0.3,actions=output:5
ubuntu@ubuntu1604:~$ sudo ovs-ofctl -O OpenFlow13 add-flow s6 in_port=1,priority
=40000,dl_type=0x800,nw_src=10.0.0.1,nw_dst=10.0.0.4,actions=output:6
ubuntu@ubuntu1604:~$

```

example, all hosts only have 1 port, and they are all connected to port 0. Host 7 is connected to port 0 through switch 6 on port 1. The below switches show different ports since they have access to multiple.

Within the network topology (Figure 1), the load balancing will take place at switch 6, port 1, on host 3 port 5, and on host 4 port 6 as shown in Figure 17. From figure 16, you can see where switch 6 is connected to port 1, and where each of the different ports are connected from the above commands. For example, host 3 is connected to port 0, through switch 5 on port 4. This load balancing is taking place through switch 1, where switch 6 uses port 5, and on switch 2, uses port 6.

Installing the flow entries means we can see the load balancing take place through the number of packets going through each switch and each port.



```
ubuntu@ubuntu1604:~$ sudo ovs-ofctl -O OpenFlow13 add-flow s6 in_port=1,priority=40000,dl_type=0x800,nw_src=10.0.0.1,nw_dst=10.0.0.3,actions=output:5
ubuntu@ubuntu1604:~$ sudo ovs-ofctl -O OpenFlow13 add-flow s6 in_port=1,priority=40000,dl_type=0x800,nw_src=10.0.0.1,nw_dst=10.0.0.4,actions=output:6
ubuntu@ubuntu1604:~$ clear
ubuntu@ubuntu1604:~$ sudo ovs-ofctl dump-flows s6 -O OpenFlow13
DPST_FLOW reply (OF1.3) (xid=0x2):
 cookie=0x0, duration=162.804s, table=0, n_packets=0, n_bytes=0, priority=40000,ip,in_port=1,nw_src=10.0.0.1,nw_dst=10.0.0.3 actions=output:5
 cookie=0x0, duration=148.554s, table=0, n_packets=0, n_bytes=0, priority=40000,ip,in_port=1,nw_src=10.0.0.1,nw_dst=10.0.0.4 actions=output:6
 cookie=0x0, duration=4530.275s, table=0, n_packets=1131, n_bytes=107062, priority=0 actions=CONTROLLER:65535
ubuntu@ubuntu1604:~$ sudo ovs-ofctl dump-flows s6 -O OpenFlow13
DPST_FLOW reply (OF1.3) (xid=0x2):
 cookie=0x0, duration=268.554s, table=0, n_packets=0, n_bytes=0, priority=40000,ip,in_port=1,nw_src=10.0.0.1,nw_dst=10.0.0.3 actions=output:5
 cookie=0x0, duration=254.304s, table=0, n_packets=0, n_bytes=0, priority=40000,ip,in_port=1,nw_src=10.0.0.1,nw_dst=10.0.0.4 actions=output:6
 cookie=0x20000154000000, duration=14.725s, table=0, n_packets=13, n_bytes=1274, idle_timeout=5, priority=1,ip,in_port=2,dl_src=a:1a:82:98:94:c0,dl_dst=26:2
5:22:2e:a7:b1,nw_src=10.0.0.1,nw_dst=10.0.0.3 actions=output:5
 cookie=0x20000155000000, duration=14.714s, table=0, n_packets=14, n_bytes=1372, idle_timeout=5, priority=1,ip,in_port=5,dl_src=26:25:22:2e:a7:b1,dl_dst=1a:b
a:82:98:94:c0,nw_src=10.0.0.3,nw_dst=10.0.0.1 actions=output:2
 cookie=0x0, duration=4636.025s, table=0, n_packets=1148, n_bytes=108335, priority=0 actions=CONTROLLER:65535
ubuntu@ubuntu1604:~$
```

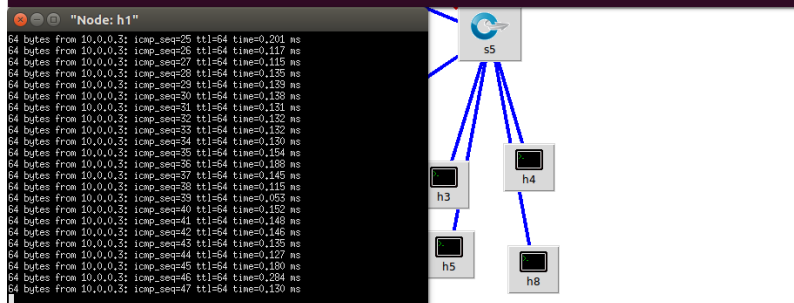


Figure 18: Testing Packet Count

Figure 19: Packet Count Increasing

Now the flow entries are installed, the packet count can be traced through the switches and hosts to ensure load balancing is occurring. Using the command line node, host 1 is opened to ping host 3 continually. The Openflow dump

```

ubuntu@ubuntu1604:~$
DPST FLOW reply (OF1.3) (xid=0x2):
cookie=0x0, duration=162.804s, table=0, n_packets=0, n_bytes=0, priority=40000,ip,in_port=1,nw_src=10.0.0.1,nw_dst=10.0.0.3 actions=output:5
cookie=0x0, duration=148.554s, table=0, n_packets=0, n_bytes=0, priority=40000,ip,in_port=1,nw_src=10.0.0.1,nw_dst=10.0.0.4 actions=output:6
cookie=0x0, duration=4530.275s, table=0, n_packets=1131, n_bytes=107602, priority=0 actions=CONTROLLER:65535
ubuntu@ubuntu1604:~$ sudo ovs-ofctl dump-flows s6 -O OpenFlow13
DPST FLOW reply (OF1.3) (xid=0x2):
cookie=0x0, duration=268.554s, table=0, n_packets=0, n_bytes=0, priority=40000,ip,in_port=1,nw_src=10.0.0.1,nw_dst=10.0.0.3 actions=output:5
cookie=0x0, duration=254.304s, table=0, n_packets=0, n_bytes=0, priority=40000,ip,in_port=1,nw_src=10.0.0.1,nw_dst=10.0.0.4 actions=output:6
cookie=0x20000154000000, duration=14.725s, table=0, n_packets=13, n_bytes=1274, idle_timeout=5, priority=1,ip,in_port=2,dl_src=1a:ba:82:98:94:c0,dl_dst=26:2
:22:2e:a7:b1,nw_src=10.0.0.1,nw_dst=10.0.0.3 actions=output:5
cookie=0x20000155000000, duration=14.744s, table=0, n_packets=14, n_bytes=1372, idle_timeout=5, priority=1,ip,in_port=5,dl_src=26:25:22:2e:a7:b1,dl_dst=1a:b
a:82:98:94:c0,nw_src=10.0.0.3,nw_dst=10.0.0.1 actions=output:2
cookie=0x0, duration=4636.025s, table=0, n_packets=1148, n_bytes=108335, priority=0 actions=CONTROLLER:65535
ubuntu@ubuntu1604:~$ sudo ovs-ofctl dump-flows s6 -O OpenFlow13
DPST FLOW reply (OF1.3) (xid=0x2):
cookie=0x0, duration=360.797s, table=0, n_packets=0, n_bytes=0, priority=40000,ip,in_port=1,nw_src=10.0.0.1,nw_dst=10.0.0.3 actions=output:5
cookie=0x0, duration=346.547s, table=0, n_packets=0, n_bytes=0, priority=40000,ip,in_port=1,nw_src=10.0.0.1,nw_dst=10.0.0.4 actions=output:6
cookie=0x20000154000000, duration=106.968s, table=0, n_packets=104, n_bytes=10192, idle_timeout=5, priority=1,ip,in_port=2,dl_src=1a:ba:82:98:94:c0,dl_dst=2
6:25:22:2e:a7:b1,nw_src=10.0.0.1,nw_dst=10.0.0.3 actions=output:5
cookie=0x20000155000000, duration=106.957s, table=0, n_packets=105, n_bytes=10290, idle_timeout=5, priority=1,ip,in_port=5,dl_src=26:25:22:2e:a7:b1,dl_dst=1
a:ba:82:98:94:c0,nw_src=10.0.0.3,nw_dst=10.0.0.1 actions=output:2
cookie=0x0, duration=4728.208s, table=0, n_packets=1164, n_bytes=109431, priority=0 actions=CONTROLLER:65535
ubuntu@ubuntu1604:~$

```

```

graph TD
    s5((s5)) --- h3[h3]
    s5 --- h4[h4]
    s5 --- h5[h5]
    s5 --- h6[h6]
    s5 --- h8[h8]

```

commands as shown in Figure 18 shows the current packet count while the ping is being carried out, and this packet count increases to 13 on host 3 and 14 on host 1.

Continuing with the packet testing, Figure 19 shows the same commands being used from host 1 to host 3, except packet count is increasing without the 2nd link being utilized. As a result, as long as only one packet counter is increasing, load balancing is shown to be taking place.

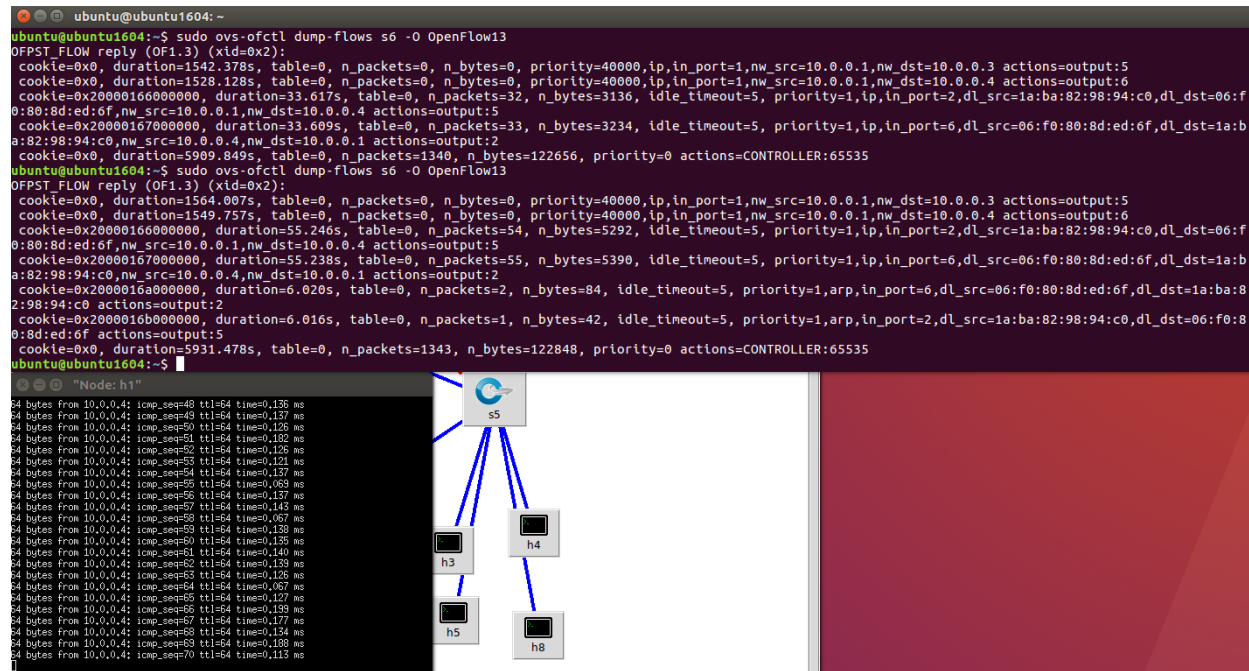
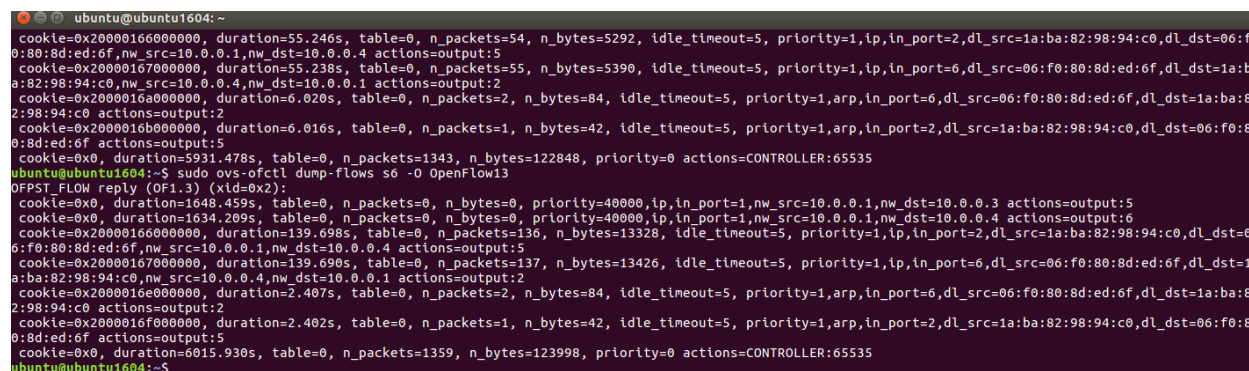


Figure 20: Packet Count from H1 to H4

Figure 21: Increasing Packet Amounts

To ensure it works correctly across all hosts, host 4



was also tested as shown in Figure 20. From the dump flow tables, packets were increasing across the switches specified for load balancing, which shows the load balancing works correctly.

Figure 21 shows load balancing working with the packet count increasing. The controller is bidirectional, meaning doing the load balancing manually on one side results in the controller automatically load balancing the other.



### 3.3. Network Monitoring

The last step of the network setup is to monitor the network to ensure all the controllers, switches and Openflow devices work correctly. Additionally, any statistics that a company want to find out can be viewed via the Floodlight rest API.

```
ubuntu@ubuntu1604:~$ sudo apt install curl
Reading package lists... Done
Building dependency tree
Reading state information... Done
The following package was automatically installed and is no longer required:
  snapd-login-service
Use 'sudo apt autoremove' to remove it.
The following additional packages will be installed:
  libcurl3-gnutls
The following NEW packages will be installed:
  curl
The following packages will be upgraded:
  libcurl3-gnutls
1 upgraded, 1 newly installed, 0 to remove and 343 not upgraded.
Need to get 328 kB of archives.
After this operation, 352 kB of additional disk space will be used.
Do you want to continue? [Y/n] y
Get:1 http://us.archive.ubuntu.com/ubuntu xenial-updates/main amd64 libcurl3-gnu
tls amd64 7.47.0-1ubuntu2.19 [189 kB]
Get:2 http://us.archive.ubuntu.com/ubuntu xenial-updates/main amd64 curl amd64 7
```

Figure 22: Curl Installation

In order to access the rest API, curl needed to be installed (Figure 22), which is used to transfer data from different servers. Alternatively, python could be used, however, using the command line interface provide ease, rather than creating a new script.

```
ubuntu@ubuntu1604: ~
ubuntu@ubuntu1604:~$ curl -X POST http://127.0.0.1:8080/wm/statistics/config/enable/json
{"statistics-collection":"enabled"}ubuntu@ubuntu1604:~$
```

Figure 23: Enable Statistics Collection

To collect statistics, the controller has to be enabled through command in Figure 23. 127.0.0.1 is the default loopback address for the controller, and the port 8080 is the default API on the controller.



```

ubuntu@ubuntu1604:~$ curl -X GET http://127.0.0.1:8080/wm/statistics/bandwidth/00:00:00:00:00:06/1/json | python -m json.tool
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
   Dload  Upload  Total    Spent    Left  Speed
100    180     0    180     0     0    1660      0 --:--:-- --:--:-- --:--:-- 1747
[
  {
    "bits-per-second-rx": "0",
    "bits-per-second-tx": "126",
    "dpid": "00:00:00:00:00:06",
    "link-speed-bits-per-second": "10000000",
    "port": "1",
    "updated": "Sun Apr 03 11:12:12 EDT 2022"
  }
]
ubuntu@ubuntu1604:~$

```

Figure 24: Bandwidth Statistics

```

ubuntu@ubuntu1604:~$ curl -X GET http://127.0.0.1:8080/wm/performance/data/json | python -m json.tool
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
   Dload  Upload  Total    Spent    Left  Speed
100   1323     0   1323     0     0   13292      0 --:--:-- --:--:-- --:--:-- 13363
{
  "average": 0,
  "current-time": "2022-04-03 11:18:16.863",
  "max": 0,
  "min": 0,
  "modules": [
    {
      "average": 0,
      "average-squared": 0,
      "max": -9223372036854775808,
      "min": 9223372036854775807,
      "module-name": "net.floodlightcontroller.firewall.Firewall",
      "num-packets": 0,
      "std-dev": 0,
      "total": 0
    },
    {
      "average": 0,
      "average-squared": 0,
      "max": -9223372036854775808,
      "min": 9223372036854775807,
      "module-name": "net.floodlightcontroller.linkdiscovery.internal.LinkDiscoveryManager",
      "num-packets": 0,
      "std-dev": 0,
      "total": 0
    },
    {
      "average": 0,
      "average-squared": 0,
      "max": -9223372036854775808,
      "min": 9223372036854775807,
      "module-name": "net.floodlightcontroller.loadbalancer.LoadBalancer",
      "num-packets": 0,
      "std-dev": 0,
      "total": 0
    },
    {
      "average": 0,
      "average-squared": 0,
      "max": -9223372036854775808,
      "min": 9223372036854775807,
      "module-name": "net.floodlightcontroller.devicemanager.internal.DeviceManagerImpl",
      "num-packets": 0,
      "std-dev": 0,
      "total": 0
    },
    {
      "average": 0,
      "average-squared": 0,

```

Figure 25: Controller performance

The first statistic is bandwidth, which is important to understand the speeds of a network. The command was a part of the Floodlight rest API as seen above and the format is json because data results are more organised. As seen from Figure 24, the statistics display bits sent/received, DPID (switch ID), link speed and the port being queried.

Figure 26: Switch Flow Statistics

Figure 25 shows the performance of the controller in terms of average packet in processing

```
ubuntu@ubuntu1604:~$ curl -X GET http://127.0.0.1:8080/wm/core/switch/all/flow/json | python -m json.tool
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
100    2017      0    2017      0      0    1626      0  --:--:--  0:00:01 --:--:-- 1661
{
  "00:00:00:00:00:00:00:01": {
    "flows": [
      {
        "byte_count": "26271",
        "cookie": "0",
        "duration_nsec": "367000000",
        "duration_sec": "1483",
        "flags": [],
        "hard_timeout_s": "0",
        "idle_timeout_s": "0",
        "instructions": {
          "instruction_apply_actions": {
            "actions": "output=controller"
          }
        },
        "match": {},
        "packet_count": "315",
        "priority": "0",
        "table_id": "0x0",
        "version": "OF_13"
      }
    ]
  },
  "00:00:00:00:00:00:00:02": {
    "flows": [
      {
        "byte_count": "26074",
        "cookie": "0",
        "duration_nsec": "337000000",
        "duration_sec": "1483",
        "flags": [],
        "hard_timeout_s": "0",
        "idle_timeout_s": "0",
        "instructions": {
          "instruction_apply_actions": {
            "actions": "output=controller"
          }
        },
        "match": {},
        "packet_count": "312",
        "priority": "0",
        "table_id": "0x0",
        "version": "OF_13"
      }
    ]
  },
  "00:00:00:00:00:00:00:03": {
```

times, which can be significant for a company when analysing performance of the controller which controls the entire network.

Figure 26 portrays the status of all switches, including packet counts, flows and table IDs. Using this can help a company establish switch performance if any issues arise.

```

ubuntu@ubuntu1604:~$ curl -X GET http://127.0.0.1:8080/wm/core/switch/00:00:00:00:00:00:06/flow/json | python -m json.tool
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload  Total   Spent    Left     Speed
100    309    0   309    0     0   8618      0  --:--:-- --:--:-- --:--:--   8828
{
  "flows": [
    {
      "byte_count": "30502",
      "cookie": "0",
      "duration_nsec": "252000000",
      "duration_sec": "1882",
      "flags": [],
      "hard_timeout_s": "0",
      "idle_timeout_s": "0",
      "instructions": {
        "instruction_apply_actions": {
          "actions": "output=controller"
        }
      },
      "match": {},
      "packet_count": "372",
      "priority": "0",
      "table_id": "0x0",
      "version": "OF_13"
    }
  ]
}
ubuntu@ubuntu1604:~$

```

Figure 27: Individual Switch Flow Statistics

Figure 28: Controller Switches

```

ubuntu@ubuntu1604:~$ curl -X GET http://127.0.0.1:8080/wm/core/controller/switches/json | python -m json.tool
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload  Total   Spent    Left     Speed
100    787    0   787    0     0   9660      0  --:--:-- --:--:-- --:--:--  10493
[
  {
    "connectedSince": 1648997807897,
    "inetAddress": "/127.0.0.1:59818",
    "openFlowVersion": "OF_13",
    "switchDPID": "00:00:00:00:00:00:00:03"
  },
  {
    "connectedSince": 1648997807897,
    "inetAddress": "/127.0.0.1:59814",
    "openFlowVersion": "OF_13",
    "switchDPID": "00:00:00:00:00:00:00:04"
  },
  {
    "connectedSince": 1648997807897,
    "inetAddress": "/127.0.0.1:59812",
    "openFlowVersion": "OF_13",
    "switchDPID": "00:00:00:00:00:00:00:06"
  },
  {
    "connectedSince": 1648997807897,
    "inetAddress": "/127.0.0.1:59820",
    "openFlowVersion": "OF_13",
    "switchDPID": "00:00:00:00:00:00:00:02"
  },
  {
    "connectedSince": 1648997807887,
    "inetAddress": "/127.0.0.1:59822",
    "openFlowVersion": "OF_13",
    "switchDPID": "00:00:00:00:00:00:00:05"
  },
  {
    "connectedSince": 1648997807897,
    "inetAddress": "/127.0.0.1:59816",
    "openFlowVersion": "OF_13",
    "switchDPID": "00:00:00:00:00:00:00:01"
  }
]
ubuntu@ubuntu1604:~$

```

Figure 27 shows a similar command switch, however, if a company wanted individual detailed statistics on a specific switch, for example, switch 6, then the above command finds those statistics.

Since switches are an important part of the network, the statistic as shown in Figure 28 portrays all of the switch DPIDs which are connected to the controller.

```
ubuntu@ubuntu1604:~$ curl -X GET http://127.0.0.1:8080/wm/core/controller/summary/json | python -m json.tool
% Total    % Received % Xferd  Average Speed   Time    Time     Time  Current
           Dload  Upload   Total     Spent    Left  Speed
100    78    0    78    0    0    2258      0  --:--:--  --:--:--  --:--:-- 13000
{
  "# Switches": 6,
  "# hosts": 12,
  "# inter-switch links": 12,
  "# quarantine ports": 0
}
```

Figure 29: Controller Link Summary

The final controller API statistic shows an overall network overview as shown in Figure 29, where all switches/hosts can be seen as an overview. This means that any new added network devices can be seen through one statistic command.